

Static Security Analysis of AWS IAM Policies for AI Workload Identities

Detecting Excessive Privilege and Identity-Pivot Risks Before Deployment

github.com/poojakira/aws-agent-identity-guard - public, default branch main.

EVIDENCE POSTURE
Verified, scoped, reproducible

Claims are constrained to the repository evidence snapshot and stated limitations.

243

Tests collected

240

Tests passed

3


Tests skipped

25

Deterministic rule IDs

SYSTEM VIEW

Architecture and Threat-Control Flow



Research Contribution

A deterministic, non-executing IAM linter specialized for AI-agent roles with 25 rule IDs spanning identity policies, trust policies, and permission-boundary risks, emitting text, JSON, and SARIF for pre-deployment review and CI gating.

Verified Evidence

EV-01	243 collected: 240 passed, 3 skipped, 0 failures/errors.	EV-02	25 deterministic rule IDs remain implemented and test-covered.
EV-03	Current performance gate over 500 synthetic policies: 0.8397 ms p95 and 2,573 policies/sec, both passing configured gates.	EV-04	Ruff/format and security-audit jobs complete successfully.
EV-05	Synthetic-policy p95: 0.8397 ms	EV-06	Synthetic-policy throughput: 2,573 policies/sec

Method and Control Path

01	Parse one IAM policy document without executing cloud actions.
02	Classify statements and trust-policy structure.
03	Apply deterministic rule checks for privilege escalation, excessive privilege, trust weaknesses, audit tampering, and missing boundaries.
04	Emit findings with stable rule IDs and severity.
05	Return CI-friendly exit codes and optional SARIF.

Limitations and Residual Risk

- No additional limitation text was found in the research brief. Consult the README before generalizing results.

Reproducibility

Expected test evidence: 240 passed, 3 skipped.

Evidence Provenance

Claim discipline Metrics come from the repository's verified metrics and research brief.
Scope discipline Benchmarks are reported only at their documented fixture, dataset, or test scope.
Reproduction PDF and PNG are generated locally from committed evidence without paid services.